

POLARIS

Performance analysis and Optimization of LARge
Infrastructures and Systems

Joint Inria / CNRS / Univ. Grenoble Alpes Project Team Proposal

Arnaud Legrand

Grenoble, May 22 2015

Project-Team Composition


- *Natural evolution* of the MESCAL team.

Name	Affiliation	Provenance	Expertise
V. Danjean	MdC UJF	MOAIS	HPC, Tracing, Experimental Methodology
N. Gast	CR2 Inria	MESCAL	Optimization, Stochastic Modeling
B. Gaujal	DR1 Inria	MESCAL	Modeling, Optimization, Game Theory
G. Huard	MdC UJF	MOAIS	HPC, Tracing, Visualization
A. Legrand	CR1 CNRS	MESCAL	HPC, Simulation, Visualization, Optimization
F. Perronnin	MdC UJF	MESCAL	Simulation, Stochastic and fluid models
P. Mertikopoulos	CR2 CNRS	MESCAL	Optimization, Game/Information Theory
J.M. Vincent	MdC UJF	MESCAL	HPC, Modeling, Simulation, Visualization

- *Inria field / theme:*
 - Network, Systems and Services
 - Distributed Computing / Distributed and High Performance Computing
- *Keywords:* HPC/large distributed systems, performance analysis, distributed and stochastic optimization, . . .

Context and Objectives

- Large distributed infrastructures
 - HPC/cloud/...
 - Wireless networks



- Common questions scalability, resilience, adaptability, capacity planning, energy consumption, ...
- Common characteristics ever growing size, distributed, heterogeneous, user-centric \leadsto stochastic nature
- This requires involved tools and new techniques that will be useful to the D&HPC community

Scientific Foundations: POLARIS in a Nutshell

Contribute to the understanding (from the observation, modeling and analysis to the optimization through adapted algorithms) of performances of very large scale distributed computing systems by applying original ideas from other research fields and application domains.

POLARIS = Team of people with the right spectrum of **skills**


Experiment design measuring/monitoring/tracing tools, experimental methodology (design, control, reproducibility)

Modeling and Simulation discrete event simulation, emulation, Markov chains, perfect sampling, Monte Carlo methods, . . .

Visualization and Statistical Analysis workload characterization (failures, parallel systems), visualization and analysis of parallel applications


Optimization stochastic approximations, mean field limits, game theory, mean field games, primal dual optimization, learning, information theory

A continuum of 5 research areas



- **Measurement** design of experiments, observation overhead control, reproducible research
- **Visualization** performance qualification and debugging, multi-scale visualization, trace comparison
- **Simulation** faithful simulation of HPC systems, sensibility/robustness, trajectory coupling
- **Fluid Modeling** local interactions, transient analysis
- **Optimization** learning algorithms in continuous nonlinear games, online and distributed optimization

A continuum of 5 research areas



- **Measurement** design of experiments, observation overhead control, reproducible research
- **Visualization** performance qualification and debugging, multi-scale visualization, trace comparison
- **Simulation** faithful simulation of HPC systems, sensibility/robustness, trajectory coupling
- **Fluid Modeling** local interactions, transient analysis
- **Optimization** learning algorithms in continuous nonlinear games, online and distributed optimization

Measurement: Reproducible Experimental Methodology

Real experiments are **costly**, **difficult** to **control** and to **reproduce**


- Cannot be studied anymore like artificial systems. Need to **inspire from other experimental fields**

Research directions:

- **Design of experiments**: involved statistical technique widely used in all fields where experiments are expensive but CS
 - **Bridge this gap** and **favor its adoption** in the D&HPC theme
- **Monitoring and tracing**: need for multi-scale (application/space/time) observation where intrusiveness is controlled
 - Evaluate the **observation/analysis quality trade-off**
- **Open science and reproducibility**: complexity and rapid technological evolution = excuse for not taking care of results reproducibility
 - Monitor/document the whole process (design, execution, data gathering, filtering, analysis)
 - Investigate/design **pragmatic workflows** to alleviate this flaw

Visualization: "Performance Driver" Identification


Traditional approach: display everything




Visualization: "Performance Driver" Identification

Traditional approach: display **everything**

→ harmful **biases** (*more information than what fits on your screen*)



Evince



Acroread

→ *overenthusiastic* use of *clustering*, *pattern mining*, *sequence alignment*

Research Directions:

- Performance **qualification** and **debugging**
 - Colleagues from D&HPC theme in deep need of new approaches/tools
- **Multi-scale** analysis (space/time/application) resilient to **noise**
 - **Entropy-based Aggregation** applied to embedded/HPC systems
- Trace **comparison**

Simulation: Very Large Stochastic Systems

- Simulation circumvents some of the previous experimental issues
 - cost/screening, extrapolation, capacity planning, ...
- Traditional approach: simplistic models to study large-scale systems, developed by D&HPC experts who know little about simulation
 - Short-lived tools with no intent of predicting anything. At best grossly indicates trend but no more expectation

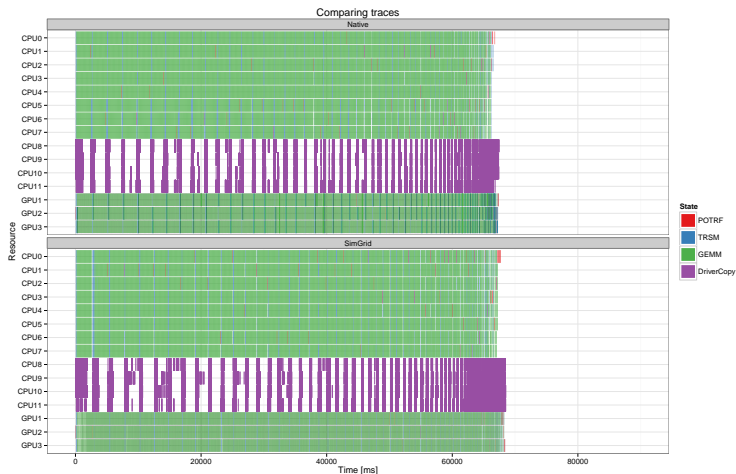
Research directions:

- Accurately reproduce the dynamic of real systems
 - SimGrid: Versatile simulation of large-scale distributed systems coarse-grain fluid models, mix emulation/simulation, invalidation
 - Used is RUNTIME/HIEPACS, ASCOLA, KERDATA, AVALON, ...
- Provide sensibility analysis and robustness indicators
- Trajectory coupling for discrete event simulations
 - PSI²: Perfect sampling for Markovian systems

Simulation: Very Large Stochastic Systems

- Simulation circumvents some of the previous experimental issues

Simulation of Cholesky/StarPU on a hybrid platform



Analysis: Local Interactions and Transient Analysis

Analysis of **stochastic** systems is particularly difficult but **mean-field** approximation is suited to **large systems**

- **Key hypothesis:** the dynamic solely depends on the entity state (not on their identity nor on their spatial location) and state space does not scale

Research directions:

- **Locality is essential:** possible approaches
 - pair approximation from statistical physics
 - fixed interaction graphs and a multi-scale approach
 - *never used for distributed computing systems and high potential*
- **Transient behavior:**
 - Finite horizon: OK (discrete system is uniformly close to its continuous limit)
 - Infinite time horizon when the continuous limit is globally stable: OK
 - Trajectory dependent stopping time: ???
 - *Could be used to analyze the complexity of distributed algorithms*

Optimization: Game Theory, On-line Distributed Optimization

Modeling interactions through game theory

Nash equilibrium often inefficient but **efficient equilibrium** can be **learned**

- Finite set of strategies = OK. **Infinite set** = ???
 - Examples: routing packet flows, power control in wireless networks, ...
 - Discretizing is not a viable option (state space explosion exponentially hard to analyze, mixed strategy space is irrelevant)
- **Goal:** Design learning algorithms in continuous nonlinear games that can be applied to realistic network scenarios

Online and distributed optimization

- Common unsatisfactory use of greedy approaches based on offline heuristics
- Each agent is faced with an **unknown and evolving loss function** and seeks to minimize his cumulative loss via the **use of past observations**
- **Regret minimization:** notion at the interface of game theory, optimization, statistics and theoretical computer science
- **Goal:** Develop and apply such techniques to actual systems

Ensure that key **practical properties are met** (asynchronous operations, numerical stability, robustness to noisy or delayed inputs, low overhead)

Distributed and H.P. Computing/Distributed Systems and Middleware
Potential or *ongoing* collaborations with: DATA-MOVE, (CORSE),
AVALON, ROMA, STORM, HIEPACS, (REALOPT), TADAAM, KER-
DATA, MYRIADS, ASAP, REGAL

Other Inria themes

- Optimization of and control of dynamic systems: BIBOP, NECS
- Networks and Telecommunications: MAESTRO, DIOGENE, DIONYSOS, RAP, SOCRATE

Other groups Game theory (LSS/supelec, Ceremade/Dauphine, HEC),
Stochastic optimization (Toulouse)

International collaborations

- Inria JLESC (NCSA/UIUC, BSC, Jülich)
- Inria@SiliconValley/Berkeley (BOINC)
- LICA (UFRGS)
- EPFL
- Univ. of Athens

Connexion with Grenoble industry through CIFRE contracts

- Bull/ATOS, STMicroelectronics, HP, Orange, CEA
- Alcatel, Huawei